

Conceptual Design. The Entity-Relationship (ER) Model

CS430/630
Lecture 12

Database Design Overview

- ▶ **Conceptual design**
 - ▶ The Entity-Relationship (ER) Model, UML
 - ▶ High-level, close to human thinking
 - ▶ Semantic model, intuitive, rich constructs
 - ▶ Not directly implementable
- ▶ **Logical Design**
 - ▶ The relational data model
 - ▶ Machine-implementable, fewer and more basic constructs
 - ▶ Logical design translates ER into relational model (SQL)
- ▶ **Physical Design** (not in this course)
 - ▶ Storage and indexing details

Conceptual Design – ER Model

- ▶ What are the *entities* and *relationships* in a typical application?
 - ▶ What information about these entities and relationships should we store in the database?
- ▶ What are the *integrity constraints* or *business rules*
 - ▶ Key constraints
 - ▶ Participation constraints
- ▶ Representation through *ER diagrams*
 - ▶ ER diagrams are then mapped into relational schemas
 - ▶ Conversion is fairly mechanical


Entities and Entity Sets

- ▶ Entity: represents a real-world object
 - ▶ Characterized using set of attributes
 - ▶ Each attribute has a domain – similar to variable types
- ▶ Entity Set: represents collection of similar entities
 - ▶ E.g., all employees in an organization
 - ▶ All entities in an entity set share same set of attributes

Keys

- ▶ Each entity set has a key
 - ▶ Set of attributes that uniquely identify an entity
 - ▶ Multiple candidate keys may exist
 - ▶ Primary key selected among them

Entity Set Representation




Representation Convention:

- Entity sets: **rectangles**
- Attributes: **ovals**, with key attributes underlined
- Edges connect entity sets to attributes

Relationships and Relationship Sets

- ▶ **Relationship**: Association among two (or more) entities
 - ▶ “Gabriel works in CS department”
 - ▶ Can have descriptive attributes: e.g., “since 9/1/2011”
 - ▶ But relationship must be fully determined by entities!
 - ▶ *Binary, ternary or multi-way (n-way) relationships*
- ▶ **Relationship Set**: Collection of similar relationships
 - ▶ Contains n -tuples (e_1, \dots, e_n) , where e_i belongs to entity set E_i
 - ▶ **Instance**: “snapshot” of relationship set at some point in time


Visualizing Relationships and Rel. Sets



Edge = Relationship

Set of Edges = Relationship Set

Relationship Set Representation

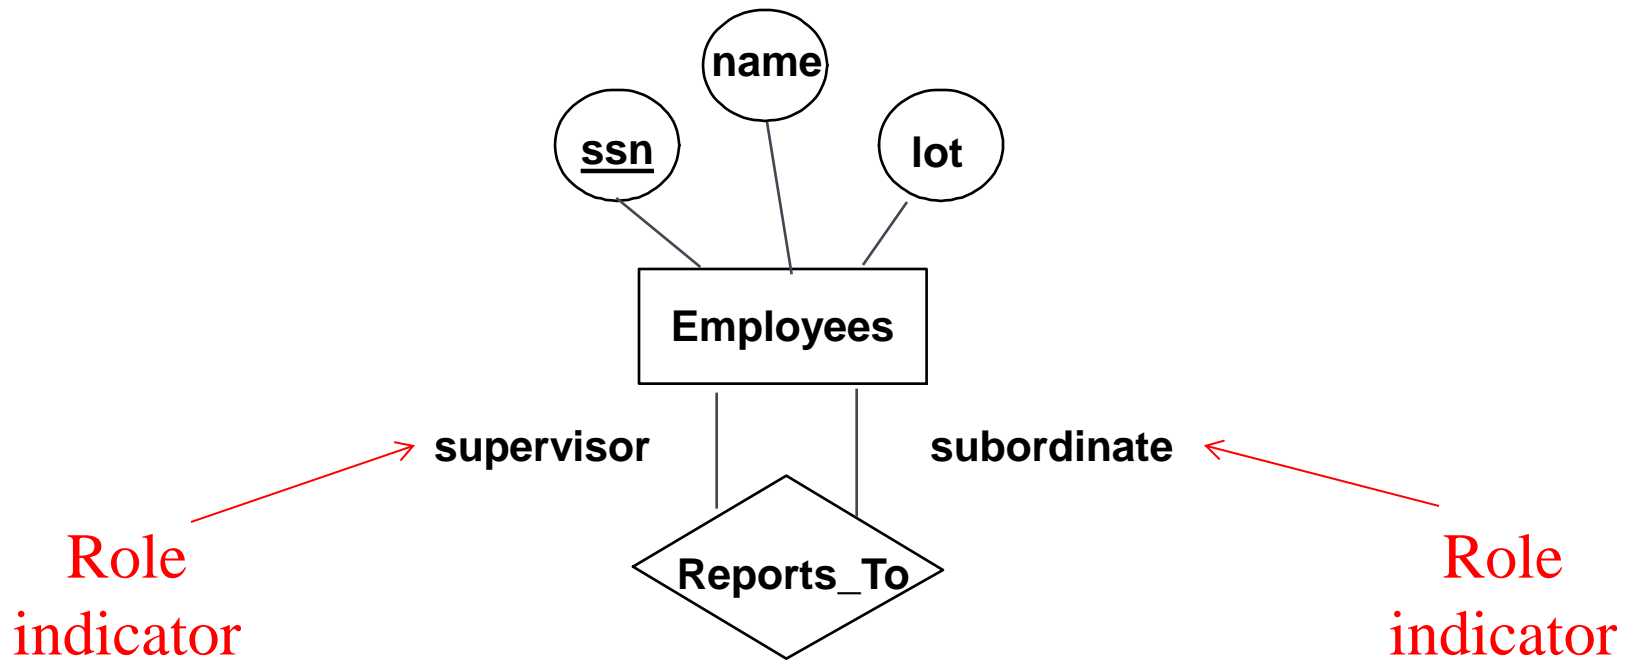


Representation Convention:

- Relationship sets: **diamonds**
- Edges connect relationship sets to entity sets, and relationship sets to relationship set attributes


A Special Case of Relationship

- ▶ An entity set can participate in a relationship set with itself
 - ▶ Entities in same set play different **roles** in the relationship
 - ▶ **Role indicators** express the role




Key Constraints


- ▶ How many other entities can an entity have a relationship with?
 - ▶ Also referred to as relationship *multiplicity*




1-to-1



1-to-Many



Many-to-1




Many-to-Many

Example 1

- ▶ **Works_In** relationship: an employee can work in many departments; a dept can have many employees.

many-to-many




Example 2

- **Manages** relationship: each dept has *at most one* manager
one-to-many

from *Employees* to *Departments* , or

many-to-one


from *Departments* to *Employees*




Participation Constraints

▶ *Total vs Partial Participation*

- ▶ **Total**: every department must have a manager
 - ▶ “Departments” entity set has total participation in relationship
 - ▶ Represented as thickened line (there is a key constraint as well)
- ▶ **Partial**: not every employee is a manager
 - ▶ “Employees” entity set has partial participation



Participation Constraints



Example


Design a database for a bank, including information about customers and their accounts. Information about customers includes their name, address, phone and SSN. Accounts have numbers, types (e.g., savings/checking) and balances.

1. Draw the E/R diagram for this database.
2. Modify the E/R diagram such that each customer must have at least one account.
3. Modify the E/R diagram further such that an account can have at most one customer.

Mapping ER to Relational Schemas

- ▶ For most part, process is mechanical
 - ▶ Some special cases arise in the presence of constraints
- ▶ Translation from ER to SQL requires:
 - ▶ Mapping entity sets to tables
 - ▶ Mapping relationship sets to tables
 - ▶ Capturing key constraints
 - ▶ Capturing participation constraints

Entity Sets to Tables




```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```

Relationship Sets to Tables


- ▶ “No-constraints” case follows simple rules
- ▶ Relationship set becomes a relation, attributes include:
 - ▶ Keys for each participating entity set (as foreign keys pointing to respective entity table)
 - ▶ All descriptive attributes for relationship
 - ▶ Primary key of relationship set table is the concatenation of primary keys for the entity sets

Relationship Sets to Tables



```
CREATE TABLE Works_In(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (ssn, did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (did)  
    REFERENCES Departments)
```

What if there are Key Constraints?



- ▶ Each department has at most one manager, according to the key constraint on Manages

Variant 1

- ▶ Map relationship to a table:
 - ▶ Note that **did** is the key now!
 - ▶ Separate table for Manages relationship.

```
CREATE TABLE Manages(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    FOREIGN KEY (did) REFERENCES Departments)
```

Variant 2


- ▶ Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget INTEGER,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```

Review: Participation Constraints

- ▶ Does every department have a manager?
 - ▶ If yes, the participation of Departments in Manages is *total*
 - ▶ Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!), but this cannot be controlled in SQL (unless we use complex constraints)
- ▶ Turns out that it is **NOT** possible to capture this with the two-tables mapping
 - ▶ Foreign key mechanism does not allow to check if there is a reference to every tuple in the referenced table
 - ▶ The Dept_Mgr variant is the only way!

Participation Constraints in SQL



```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget INTEGER,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees  
  ON DELETE NO ACTION)
```


Participation Constraints Summary

- ▶ **General case**
 - ▶ Total participation cannot be enforced unless we use complex constraints
- ▶ **What if there is also a key constraint in place?**
 - ▶ If the entity set with total participation also has a key constraint, then it is possible to capture total participation
 - ▶ But only if “combined” table construction is used!

Design Choices in the ER Model

- ▶ Should a concept be modeled as an entity or an attribute?
- ▶ Should a concept be modeled as an entity or a relationship?
 - ▶ Considers hierarchies and inheritance
 - ▶ Outside the scope of this class

Entity vs. Attribute



- ▶ Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?

Entity vs. Attribute

- ▶ Sometimes **address** may have to be an entity:
 - ▶ If we have several addresses per employee (since attributes cannot be set-valued)
 - ▶ If the structure (city, street, etc.) is important, e.g., retrieve employees in a given city (attribute values are atomic!)

